

23 — Company-Specific Interview Preparation

Company-specific PostgreSQL and SQL interview guides for top tech companies.

Files

File	Company	Key Focus
01_amazon_preparation.md	Amazon / AWS	LP principles, Aurora PostgreSQL, order management design
02_google_preparation.md	Google	Query optimization, distributed systems (Spanner), BigQuery
03_microsoft_preparation.md	Microsoft	Azure PostgreSQL, T-SQL vs PG differences, enterprise features
04_meta_preparation.md	Meta	Petabyte-scale SQL, Presto, social graph, A/B testing
05_uber_preparation.md	Uber	PostGIS geospatial (15 examples), H3 hexagons, real-time location
06_netflix_preparation.md	Netflix	Streaming analytics, A/B testing, recommendation systems
07_general_tech_companies.md	Multi-company	Airbnb, Stripe, DoorDash, Snowflake, Databricks, LinkedIn, Atlassian

Each File Contains

1. Company Profile (culture, DB tech stack, what they value)
2. Interview Process (rounds, format, duration, level expectations)
3. SQL Coding Round (20–25 questions with full solutions)
4. Database Design Tasks (3–5 complete schema designs)
5. System Design with DB Focus (2–3 scenarios)
6. Behavioral Questions (role-specific)
7. Mock Interview Script (full Q&A dialogue)
8. Evaluation Rubric (L3/L4/L5/L6 bar)
9. Red Flags (what tanks candidates)
10. Green Flags (what impresses interviewers)
11. Salary Bands (approximate total comp by level)
12. Preparation Timeline (4-week plan)

Quick Start by Company

If interviewing at Amazon

- Focus on Leadership Principles first
- Study Aurora PostgreSQL vs. vanilla PG differences
- Practice LP storytelling with DB metrics

If interviewing at Google

- Master EXPLAIN ANALYZE and execution plans

- Study Cloud Spanner and BigQuery alongside PostgreSQL
- Practice query optimization patterns

If interviewing at Microsoft

- Learn Azure PostgreSQL Flexible Server specifics
- Study row-level security, encryption, compliance features
- Practice Azure-flavored system design

If interviewing at Meta

- Every design must start at "billions of rows"
- Study Presto/Trino SQL for analytics
- Practice social graph SQL patterns

If interviewing at Uber

- PostGIS is mandatory — practice all 15 geospatial queries
- Understand H3 hexagonal grid indexing
- Practice real-time vs. durable data split patterns

If interviewing at Netflix

- Study streaming analytics patterns (Kafka, Flink)
- Practice A/B test analysis SQL
- Understand Freedom & Responsibility culture

If interviewing at Stripe, DoorDash, Atlassian, etc.

- Study the `07_general_tech_companies.md` company-specific sections
- Match SQL style to company domain
- Research the engineering blog before the interview

Common SQL Patterns Tested Everywhere

```
-- Top N per group
SELECT * FROM (...RANK() OVER (PARTITION BY cat ORDER BY val DESC)... ) WHERE rnk <=
N;

-- Cohort retention
-- Date spine with generate_series + LEFT JOIN for gap filling
-- Running totals and moving averages
-- Session analysis (gap-based grouping)
-- Funnel conversion rates
-- Year-over-year comparison with LAG or date filtering
```

Salary Ranges Summary (USD, 2024–2025, mid-level/senior)

Company	Mid-Level Total Comp	Senior Total Comp
Google L5	\$360K–\$570K	\$560K–\$900K
Meta E5	\$390K–\$620K	\$700K–\$1.2M
Netflix L5	\$310K–\$440K	\$440K–\$700K

Amazon SDE III	\$290K-\$420K	\$450K-\$700K
Microsoft Sr SDE	\$280K-\$420K	\$440K-\$700K
Stripe L3	\$300K-\$450K	\$450K-\$650K
Snowflake Senior	\$380K-\$600K	\$600K-\$950K
Uber L5	\$300K-\$480K	\$500K-\$750K
Databricks Senior	\$400K-\$650K	\$650K-\$950K