

Learning Paths — PostgreSQL Complete Guide

Choose your path based on your current role, experience level, and goals. Each path is designed to be efficient — you will cover what matters most for your context without wasting time on irrelevant material.

Table of Contents

1. [How to Choose Your Path](#)
 2. [Path A — Complete Beginner](#)
 3. [Path B — Backend Engineer](#)
 4. [Path C — Data Engineer](#)
 5. [Path D — Analytics Engineer](#)
 6. [Path E — Database Engineer / DBA](#)
 7. [Path F — Interview-Only Fast Track](#)
 8. [Cross-Path Reference Table](#)
 9. [Chapter Dependency Map](#)
 10. [Milestone Definitions](#)
-

1. How to Choose Your Path

Answer these questions:

Q1: Have you written SQL before?

- No → Path A (Complete Beginner)
- Yes, basics only → Go to Q2

Q2: What is your primary goal?

- Build production backend services with PostgreSQL → Path B
- Build data pipelines and ETL workflows → Path C
- Build analytics queries and reports → Path D
- Administer, tune, and operate PostgreSQL → Path E
- Pass an interview in the next 3–4 weeks → Path F

Q3: How much time per week can you commit?

- 5–10 hours/week → Double the estimated durations
 - 10–15 hours/week → Use the given estimates
 - 15–20+ hours/week → Reduce estimates by 25%
-

2. Path A — Complete Beginner

Target audience: No prior SQL or database experience. Starting from absolute zero.

Estimated duration: 16–20 weeks at 10 hours/week

End goal: Junior-to-mid level PostgreSQL proficiency. Ready for junior backend, junior data, or junior analyst roles.

Phase 1 — Foundations (Weeks 1–3)

Learning objectives:

- Understand what relational databases are and why they exist
- Write basic SELECT queries with filtering and sorting
- Understand the client/server model and connect to PostgreSQL

Week	Chapter	Topic	Hours
1	01_Fundamentals	Relational model, SQL history, psql basics	10h
2	02_SQL_Basics (Part 1)	SELECT, WHERE, ORDER BY, LIMIT, basic functions	10h
3	02_SQL_Basics (Part 2)	JOINS, GROUP BY, HAVING, aggregates	10h

Milestone 1 checkpoint: Can you write a query that joins three tables, filters by date range, groups by category, and returns the top 5 results? If yes, proceed.

Mini-project: Query a sample e-commerce database (provided in `22_Projects/project_01`) to answer 10 business questions.

Phase 2 — Intermediate SQL (Weeks 4–6)

Learning objectives:

- Use subqueries and CTEs to break down complex problems
- Apply window functions for ranking and running totals
- Handle NULL values correctly

Week	Chapter	Topic	Hours
4	03_Intermediate_SQL (Part 1)	Subqueries (scalar, correlated), CTEs	10h
5	03_Intermediate_SQL (Part 2)	Window functions: ROW_NUMBER, RANK, LEAD, LAG	10h
6	03_Intermediate_SQL (Part 3)	Window frames, PARTITION BY, advanced patterns	10h

Milestone 2 checkpoint: Can you calculate a 7-day rolling average using window functions? Can you write a recursive CTE? If yes, proceed.

Mini-project: Build a cohort retention analysis from a user events table.

Phase 3 — Advanced SQL (Week 7)

Learning objectives:

- Solve complex analytical problems with advanced SQL patterns
- Write recursive queries for hierarchical data

Week	Chapter	Topic	Hours
7	04_Advanced_SQL	Recursive CTEs, lateral joins, gaps and islands, pivot	10h

Phase 4 — PostgreSQL Core (Weeks 8-9)

Learning objectives:

- Understand PostgreSQL's rich type system
- Use JSONB, arrays, and PostgreSQL-specific features
- Manage schemas and extensions

Week	Chapter	Topic	Hours
8	05_PostgreSQL_Core (Part 1)	Data types, type casting, arrays, JSONB	10h
9	05_PostgreSQL_Core (Part 2)	Extensions, schemas, search_path, pg_catalog	10h

Phase 5 — Database Design (Week 10)

Learning objectives:

- Design normalized schemas from business requirements
- Apply constraints correctly

Week	Chapter	Topic	Hours
10	06_Database_Design	Normalization, ERD, constraints, naming	10h

Milestone 3 checkpoint: Can you design a schema for a simple SaaS application (users, organizations, subscriptions, invoices) with proper normalization and constraints?

Phase 6 — Indexes and Optimization (Weeks 11-13)

Learning objectives:

- Understand how B-Tree indexes work
- Read basic EXPLAIN output
- Choose when to add an index

Week	Chapter	Topic	Hours
11	07_Indexes	B-Tree, partial indexes, expression indexes, GIN	10h
12	08_Query_Optimization (Part 1)	EXPLAIN, EXPLAIN ANALYZE, plan nodes	10h
13	08_Query_Optimization (Part 2)	Statistics, planner settings, fixing slow queries	10h

Phase 7 — Transactions and Production Basics (Weeks 14-16)

Learning objectives:

- Understand ACID and MVCC
- Know basic isolation levels
- Understand connection pooling

Week	Chapter	Topic	Hours
------	---------	-------	-------

14	09_Transactions_Concurrency	ACID, MVCC, isolation levels, locking basics	10h
15	12_Production_PostgreSQL	Config basics, PgBouncer intro, monitoring	10h
16	19_Interview_Questions (Junior section)	Interview prep, Q&A review	10h

Milestone 4 — Graduation checkpoint: Complete the Junior Interview Questions section in `19_Interview_Questions` . Score at least 80% without notes.

Recommended Resources for Path A

- Chapters 01–09 (core path)
- Chapter 19 (junior questions)
- Chapter 24 (cheat sheets — use throughout)
- Project 1 from `22_Projects`

3. Path B — Backend Engineer

Target audience: Software engineers who build APIs and services. You know programming. You may know basic SQL. You need to use PostgreSQL correctly in production systems.

Estimated duration: 8–10 weeks at 10 hours/week

Prerequisite: Basic SQL knowledge (can write SELECT, JOIN, GROUP BY). If not, do Path A Phases 1–3 first (3 weeks).

End goal: Can design production schemas, write efficient queries, manage connections, handle migrations safely, and debug production database issues.

Week 1 — PostgreSQL Core and Schema Design

Day	Topic	Chapter	Hours
1–2	PostgreSQL data types deep dive	05_PostgreSQL_Core	3h
3–4	JSONB patterns for backend engineers	05_PostgreSQL_Core, 11_Advanced_PostgreSQL	3h
5–7	Schema design for SaaS, multi-tenancy, UUIDs	06_Database_Design	4h

Focus areas: When to use JSONB vs normalized, UUID vs serial, schema naming conventions, constraint design.

Week 2 — Indexes for Production

Day	Topic	Chapter	Hours
1–2	B-Tree internals and multi-column ordering	07_Indexes	3h

3–4	GIN for JSONB and arrays, partial indexes	07_Indexes	3h
5–7	Covering indexes, index bloat, CREATE INDEX CONCURRENTLY	07_Indexes	4h

Focus areas: Never create an index without understanding why. `CREATE INDEX CONCURRENTLY` for zero-downtime.

Week 3 — Query Optimization

Day	Topic	Chapter	Hours
1–3	EXPLAIN ANALYZE — reading every node type	08_Query_Optimization	4h
4–5	Fixing N+1 in ORM-generated queries	08_Query_Optimization, 17_Backend_Engineers	3h
6–7	pg_stat_statements setup and usage	08_Query_Optimization	3h

Week 4 — Transactions, Locking, and Concurrency

Day	Topic	Chapter	Hours
1–2	ACID guarantees and when they matter	09_Transactions_Concurrency	3h
3–4	MVCC, isolation levels, read phenomena	09_Transactions_Concurrency	3h
5–7	SELECT FOR UPDATE, advisory locks, deadlock prevention	09_Transactions_Concurrency	4h

Focus areas: How to implement optimistic locking, how to use advisory locks as a distributed lock, how to avoid deadlocks in application code.

Week 5 — Backend-Specific Patterns

Day	Topic	Chapter	Hours
1–2	Connection pooling: PgBouncer transaction mode	17_Backend_Engineers	3h
3–4	Schema migrations: zero-downtime strategies	17_Backend_Engineers	3h
5–7	LISTEN/NOTIFY, upsert patterns, bulk operations	17_Backend_Engineers	4h

Week 6 — Production Operations

Day	Topic	Chapter	Hours
1–2	postgresql.conf tuning for backend workloads	12_Production_PostgreSQL	3h
3–4	Monitoring: pg_stat_*, slow query log, Grafana	12_Production_PostgreSQL	3h

5–7	Bloat monitoring, VACUUM basics, autovacuum	12_Production_PostgreSQL, 10_Internals	4h
-----	---	--	----

Week 7 — Security

Day	Topic	Chapter	Hours
1–3	Row-level security for multi-tenant apps	14_Security	4h
4–5	Roles, privileges, SECURITY DEFINER	14_Security	3h
6–7	pg_hba.conf, SSL, connection security	14_Security	3h

Week 8 — Architecture and Case Studies

Day	Topic	Chapter	Hours
1–3	Multi-tenant architecture patterns	18_Architecture_Case_Studies	4h
4–5	Backend system design with PostgreSQL	21_System_Design	3h
6–7	Interview prep for backend roles	19_Interview_Questions	3h

Backend Engineer Milestones:

- ☐ Can implement a zero-downtime migration adding an index to a 100M-row table
- ☐ Can explain why `SELECT FOR UPDATE` prevents double-booking
- ☐ Can configure PgBouncer in transaction mode
- ☐ Can read `EXPLAIN ANALYZE` output and identify the bottleneck
- ☐ Can implement row-level security for a multi-tenant application

4. Path C — Data Engineer

Target audience: Engineers who build data pipelines, ETL/ELT workflows, and data infrastructure.

Estimated duration: 8–10 weeks at 10 hours/week

Prerequisite: Intermediate SQL (can write CTEs and subqueries). Python or Spark experience helpful.

End goal: Can design and operate a PostgreSQL-based data platform with partitioning, CDC, FDWs, and efficient bulk operations.

Week 1 — Advanced SQL for Data Engineering

Topic	Chapter	Hours
Window functions for analytics	03_Intermediate_SQL	5h
Recursive CTEs for hierarchical data	04_Advanced_SQL	5h

Week 2 — PostgreSQL Data Types for Data Engineering

Topic	Chapter	Hours
Arrays, ranges, JSONB for semi-structured data	05_PostgreSQL_Core	4h
Full-text search for pipeline metadata	11_Advanced_PostgreSQL	3h
Custom types and domains	11_Advanced_PostgreSQL	3h

Week 3 — Query Optimization for Analytical Queries

Topic	Chapter	Hours
Reading EXPLAIN for large analytical queries	08_Query_Optimization	4h
Parallel query configuration	08_Query_Optimization	3h
Statistics and query planner for analytics	08_Query_Optimization	3h

Week 4 — Partitioning

Topic	Chapter	Hours
Range partitioning for time-series data	16_Data_Engineering	4h
List and hash partitioning	16_Data_Engineering	3h
Partition pruning, maintenance, and gotchas	16_Data_Engineering	3h

Milestone: Design a partitioned event table for 1 billion rows. Verify partition pruning in `EXPLAIN` .

Week 5 — Bulk Loading and ETL Patterns

Topic	Chapter	Hours
COPY command for bulk loading	16_Data_Engineering	3h
UNLOGGED tables and staging patterns	16_Data_Engineering	3h
INSERT ... ON CONFLICT for upserts	16_Data_Engineering	2h
Incremental loading strategies	16_Data_Engineering	2h

Week 6 — Change Data Capture

Topic	Chapter	Hours
Logical replication as CDC	11_Advanced_PostgreSQL	3h
Publications and subscriptions	13_Replication_HA	3h
Debezium integration with PostgreSQL	16_Data_Engineering	4h

Week 7 — Foreign Data Wrappers and Federated Queries

Topic	Chapter	Hours
postgres_fdw setup and query pushdown	11_Advanced_PostgreSQL	4h
file_fdw for CSV ingestion	11_Advanced_PostgreSQL	2h
Performance and limitations of FDWs	11_Advanced_PostgreSQL	2h
Federated query patterns	16_Data_Engineering	2h

Week 8 — Production Data Platform

Topic	Chapter	Hours
Monitoring data platform health	12_Production_PostgreSQL	3h
VACUUM and maintenance for large tables	12_Production_PostgreSQL	3h
dbt integration with PostgreSQL	16_Data_Engineering	4h

Week 9 — Internals Relevant to Data Engineering

Topic	Chapter	Hours
WAL and how it enables CDC	10_PostgreSQL_Internals	3h
TOAST and large column storage	10_PostgreSQL_Internals	2h
Bloat in high-update tables	10_PostgreSQL_Internals	2h
Checkpoint and WAL settings for bulk loads	12_Production_PostgreSQL	3h

Data Engineer Milestones:

- ☐ Can design and implement a partitioned table with automated partition creation
- ☐ Can set up logical replication as a CDC source
- ☐ Can implement a bulk load that processes 10M rows/minute with COPY
- ☐ Can query data across two PostgreSQL instances with FDW
- ☐ Can build an incremental dbt model on a partitioned table

5. Path D — Analytics Engineer

Target audience: Analytics engineers, BI developers, data analysts who write complex SQL and build data models.

Estimated duration: 6–8 weeks at 10 hours/week

Prerequisite: Basic SQL. Can write SELECT, JOIN, GROUP BY.

End goal: Can write world-class analytical SQL, understand query performance, build maintainable data models, and contribute meaningfully to data infrastructure decisions.

Week 1 — Intermediate SQL Mastery

Topic	Chapter	Hours
CTEs — composability and performance	03_Intermediate_SQL	4h
Window functions — all frame types	03_Intermediate_SQL	6h

Drill: Write 20 window function queries. Track time. Target: < 3 minutes per query.

Week 2 — Advanced SQL Patterns

Topic	Chapter	Hours
Conditional aggregation (FILTER, CASE-WHEN)	04_Advanced_SQL	3h
Pivot and unpivot	04_Advanced_SQL	2h
Gaps and islands problems	04_Advanced_SQL	2h
Date and time patterns for analytics	04_Advanced_SQL	3h

Week 3 — PostgreSQL for Analytics

Topic	Chapter	Hours
PostgreSQL-specific aggregate functions	05_PostgreSQL_Core	3h
JSONB for semi-structured event data	05_PostgreSQL_Core	3h
Full-text search for search analytics	11_Advanced_PostgreSQL	4h

Week 4 — Query Performance

Topic	Chapter	Hours
Reading EXPLAIN ANALYZE for analytical queries	08_Query_Optimization	5h
Index selection for analytics	07_Indexes	5h

Focus: Understand hash join vs merge join vs nested loop. Know when each is appropriate.

Week 5 — Data Modeling for Analytics

Topic	Chapter	Hours
Dimensional modeling in PostgreSQL	06_Database_Design	3h
Materialized views — creation and refresh	11_Advanced_PostgreSQL	3h
Partitioning for analytical workloads	16_Data_Engineering	4h

Week 6 — dbt and Analytics Tooling

Topic	Chapter	Hours
dbt models on PostgreSQL	16_Data_Engineering	4h
Incremental materialization strategies	16_Data_Engineering	3h
Monitoring query performance	12_Production_PostgreSQL	3h

Week 7 — Interview Preparation

Topic	Chapter	Hours
SQL interview question patterns	19_Interview_Questions	4h
Analytical SQL tasks	20_Interview_Tasks	6h

Analytics Engineer Milestones:

- ☐ Can write a cohort retention analysis using window functions
- ☐ Can build a date spine using `generate_series` and cross joins
- ☐ Can explain why a materialized view refresh is blocking queries
- ☐ Can read `EXPLAIN ANALYZE` and identify if an index is being used
- ☐ Can implement a slowly changing dimension in PostgreSQL

6. Path E — Database Engineer / DBA

Target audience: Database administrators, infrastructure engineers, SREs who are responsible for operating and scaling PostgreSQL in production.

Estimated duration: 12–14 weeks at 10 hours/week

Prerequisite: Comfortable with SQL. Some Linux administration experience.

End goal: Can design, deploy, operate, tune, secure, and recover production PostgreSQL systems at scale. Ready for senior/staff DBA roles.

Phase 1 — Deep SQL and PostgreSQL Fundamentals (Weeks 1–2)

Week	Topic	Chapter
1	Advanced SQL — all patterns	03, 04
2	PostgreSQL core, types, extensions	05

Phase 2 — Schema Design and Indexing (Weeks 3–4)

Week	Topic	Chapter
3	Schema design, normalization, constraints	06

4	All index types — internals and selection	07
---	---	----

Milestone: Can audit an existing schema and produce a written report of normalization issues, missing indexes, and redundant indexes.

Phase 3 — Query Optimization (Weeks 5–6)

Week	Topic	Chapter
5	Query planner, statistics, EXPLAIN deep dive	08
6	pg_stat_statements, auto_explain, slow query workflow	08, 12

Milestone: Given a slow query, produce a written optimization plan including: current plan analysis, root cause, proposed fix, before/after timing.

Phase 4 — Transactions and Internals (Weeks 7–8)

Week	Topic	Chapter
7	MVCC, isolation levels, locking, deadlocks	09
8	WAL, heap, TOAST, page structure, shared buffers	10

Milestone: Can draw the lifecycle of a single-row UPDATE from client to WAL to shared buffer to heap page to visibility map. Can explain xmin/xmax semantics.

Phase 5 — VACUUM and Autovacuum (Weeks 9)

Week	Topic	Chapter
9	VACUUM internals, autovacuum tuning, bloat management	10, 12

Milestone: Can identify a table with autovacuum lag, diagnose the cause, and implement a fix. Can calculate the autovacuum trigger threshold.

Phase 6 — Replication and HA (Weeks 10–11)

Week	Topic	Chapter
10	Streaming replication, physical standby, replication slots	13
11	Logical replication, Patroni, PgBouncer HA setup	13

Milestone: Can deploy a Patroni cluster with three nodes (primary + two standbys) using etcd as DCS. Can perform a controlled switchover without data loss.

Phase 7 — Backup, Recovery, and Security (Weeks 12–13)

Week	Topic	Chapter
12	PITR, pgBackRest, backup verification	15

13	RLS, pgaudit, roles, pg_hba.conf hardening	14
----	--	----

Phase 8 — Interview Preparation (Week 14)

Week	Topic	Chapter
14	Senior/Staff interview Q&A, system design	19, 21

Database Engineer Milestones:

- ☐ Can explain MVCC without notes
- ☐ Can read a full `EXPLAIN (ANALYZE, BUFFERS)` output
- ☐ Can configure autovacuum per-table overrides
- ☐ Can set up streaming replication from scratch
- ☐ Can perform a PITR restore and verify data integrity
- ☐ Can implement row-level security for a multi-tenant schema
- ☐ Can deploy PgBouncer in transaction mode
- ☐ Can design a Patroni HA cluster in a system design interview

7. Path F — Interview-Only Fast Track

Target audience: Engineers with existing SQL/PostgreSQL experience who need to pass an interview in 3–4 weeks.

Estimated duration: 3–4 weeks at 15 hours/week

Prerequisite: Intermediate SQL proficiency. Some PostgreSQL experience.

End goal: Pass technical interview rounds at mid-to-senior level at any target company.

Week 1 — SQL Skills Refresh and Gap-Filling

Day	Focus	Chapter	Hours/Day
1	Window functions — all patterns	03_Intermediate_SQL	3h
2	CTEs, recursive CTEs	03, 04	3h
3	Gaps and islands, pivot patterns	04_Advanced_SQL	3h
4	EXPLAIN ANALYZE — read any plan	08_Query_Optimization	3h
5	Index selection and design	07_Indexes	3h
6–7	SQL problem sets: 30 problems timed	20_Interview_Tasks	5h

Day 7 checkpoint: Complete the "Mid-Level SQL" section of `20_Interview_Tasks`. Target: solve 80% correctly without hints.

Week 2 — PostgreSQL Internals and Concepts

Day	Focus	Chapter	Hours/Day
1	MVCC, ACID, isolation levels	09_Transactions_Concurrency	3h
2	WAL, checkpoint, crash recovery	10_PostgreSQL_Internals	3h
3	VACUUM, autovacuum, bloat	10, 12	3h
4	Replication: streaming + logical	13_Replication_HA	3h
5	Backup: PITR concepts	15_Backup_Recovery	2h
6	Security: RLS, roles, pg_hba.conf	14_Security	2h
7	Practice explaining internals out loud	19_Interview_Questions	3h

Week 3 — System Design and Company Prep

Day	Focus	Chapter	Hours/Day
1	System design framework for DB-heavy systems	21_System_Design	3h
2	HA design with Patroni	18, 21	3h
3	Sharding and partitioning design	18, 21	3h
4	Company-specific prep (your target)	23_Company_Preparation	4h
5	Behavioral questions for DB engineers	INTERVIEW_PREPARATION.md	2h
6–7	Mock interview simulation (full rounds)	20_Interview_Tasks	5h

Week 4 — Final Polish

Day	Focus	Hours
1	Review all red/yellow flashcard items	2h
2	Senior question bank — all topics	3h
3	System design mock: design a HA DB for a ride-share app	2h
4	Company-specific research + prep	2h
5	Light review, cheat sheets, rest	1h
6–7	Mock interviews with a friend	4h

Fast Track Milestones:

- ☐ Can answer "Explain MVCC" in 3 minutes without notes
- ☐ Can read any EXPLAIN ANALYZE output and identify the problem
- ☐ Can write a window function query in < 2 minutes
- ☐ Can design a HA PostgreSQL system on a whiteboard in 15 minutes

- ☐ Completed 50+ timed SQL problems


8. Cross-Path Reference Table

Chapter	Beginner	Backend	Data Eng	Analytics	DBA	Fast Track
01_Fundamentals	Core	Skip	Skip	Skip	Skip	Skip
02_SQL_Basics	Core	Skip	Skip	Core	Skip	Skim
03_Intermediate_SQL	Core	Skim	Core	Core	Skim	Core
04_Advanced_SQL	Core	Optional	Core	Core	Optional	Core
05_PostgreSQL_Core	Core	Core	Core	Core	Core	Skim
06_Database_Design	Core	Core	Optional	Optional	Core	Skim
07_Indexes	Core	Core	Optional	Core	Core	Core
08_Query_Optimization	Core	Core	Core	Core	Core	Core
09_Transactions_Concurrency	Core	Core	Core	Optional	Core	Core
10_PostgreSQL_Internals	Optional	Optional	Core	Skip	Core	Core
11_Advanced_PostgreSQL	Optional	Optional	Core	Core	Core	Skim
12_Production_PostgreSQL	Optional	Core	Core	Optional	Core	Skim
13_Replication_HA	Skip	Optional	Core	Skip	Core	Core
14_Security	Skip	Core	Optional	Skip	Core	Skim
15_Backup_Recovery	Skip	Optional	Optional	Skip	Core	Skim
16_Data_Engineering	Skip	Optional	Core	Core	Optional	Skip
17_Backend_Engineers	Skip	Core	Skip	Skip	Optional	Skim
18_Architecture	Skip	Core	Optional	Optional	Core	Core
19_Interview_Questions	Core	Core	Core	Core	Core	Core
20_Interview_Tasks	Core	Core	Core	Core	Core	Core
21_System_Design	Skip	Core	Optional	Skip	Core	Core
22_Projects	Core	Core	Core	Core	Core	Skip
23_Company_Preparation	Skip	Optional	Optional	Optional	Optional	Core
24_Cheat_Sheets	All	All	All	All	All	Core
25_Capstone	Optional	Optional	Optional	Optional	Core	Skip

Legend: **Core** = must complete, **Optional** = do if time allows, **Skim** = read headings only, **Skip** = not relevant

9. Chapter Dependency Map

Some chapters have strict prerequisite chains. Do not skip prerequisites.



14_Security ← requires 05, 06, 09

17_Backend_Engineers ← requires 05, 06, 07, 08, 09

18_Architecture ← requires 06, 09, 13

21_System_Design ← requires 13, 18

22_Projects ← requires relevant preceding chapters

23_Company_Preparation ← requires 19, 20, 21

10. Milestone Definitions

These milestone definitions are used across all paths. Each milestone represents a concrete skill checkpoint.

Milestone 1 — SQL Foundations

Can: Write a multi-table JOIN with GROUP BY, filtering, and HAVING. Handle NULLs correctly. Use all aggregate functions. **Chapter:** 02_SQL_Basics **Test:** Complete the SQL Basics problem set in 20_Interview_Tasks with 80% accuracy.

Milestone 2 — SQL Practitioner

Can: Write CTEs, subqueries, and window functions. Read a basic EXPLAIN output. **Chapter:** 03_Intermediate_SQL **Test:** Complete the Intermediate SQL section of 20_Interview_Tasks without hints.

Milestone 3 — PostgreSQL Developer

Can: Design a normalized schema, choose indexes correctly, explain MVCC. **Chapters:** 05–09 **Test:** Complete the Mid-Level question bank in 19_Interview_Questions with 70% accuracy.

Milestone 4 — Senior PostgreSQL Engineer

Can: Configure a production instance, read EXPLAIN ANALYZE, tune autovacuum, implement RLS, explain WAL. **Chapters:** 10–15 **Test:** Complete the Senior question bank in 19_Interview_Questions with 70%

accuracy. Complete Project 3 or 4 in `22_Projects` .

Milestone 5 — Staff/Principal PostgreSQL Engineer

Can: Design a Patroni HA cluster, implement PITR, design a sharding strategy, lead a system design interview.

Chapters: All **Test:** Complete the Capstone in `25_Capstone_Program` . Complete a mock system design interview with a peer.